

# **SPARK**

## **T. VERBEIREN**

**9/7/2014**

# CONTENTS

Introduction

Spark

Ecosystem

Not covered

Example(s)

**ME, MYSELF  
AND I**

**PhD on Neural Networks**



**ICT & Management Consulting**



**ML and Data Management**



**Visual Analytics @ KU Leuven**





# INTRODUCTION

# DISTRIBUTION IS HARD ...

2 world views:

- **H**igh **P**erformance **C**omputing
- **H**igh **T**hroughput **C**omputing

# MAP / REDUCE

**MAP**

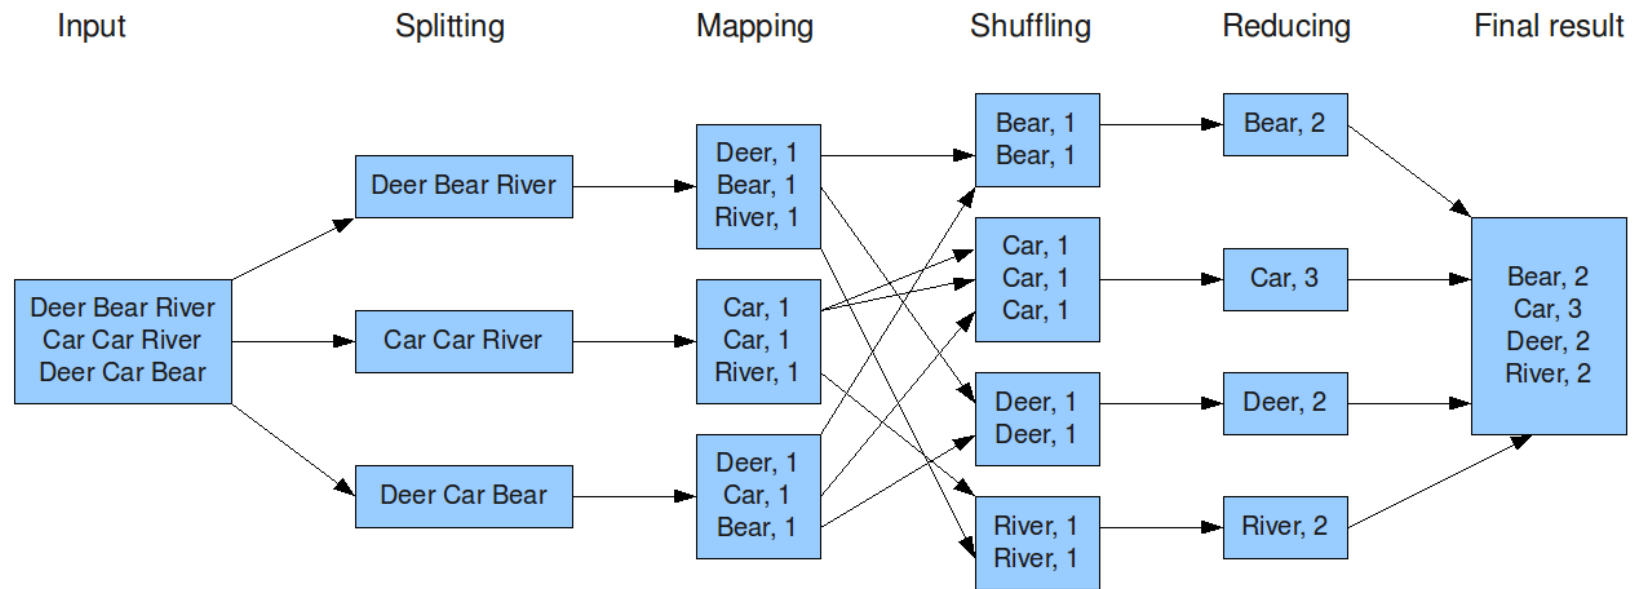
**REDUCE**

(Nothing special ?!)



# WORD COUNT IN M/R

The overall MapReduce word count process



```
import java.io.IOException;
import java.util.*;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;

public class WordCount {

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();

        Job job = new Job(conf, "wordcount");
```

**WHAT IF ...**

# ... WE COULD JUST WRITE

```
val y = x map () reduce ()
```

# ... THIS COULD BE EXTENDED

```
val y = x map () filter() map () flatMap () reduce ()
```

# WHAT WOULD BE NEEDED?

1. Platform
2. Parallel abstraction mechanism
3. Language support

# SPARK





- Berkeley University
- Apache Project
- v1.0 released on May 30d 2014
- Written in Scala
- Supported by DataBricks
- Used by ...

# LANGUAGE SUPPORT

## Language Support

### Python

```
lines = sc.textFile(...)
lines.filter(lambda s: "ERROR" in s).count()
```

### Scala

```
val lines = sc.textFile(...)
lines.filter(x => x.contains("ERROR")).count()
```

### Java

```
JavaRDD<String> lines = sc.textFile(...);
lines.filter(new Function<String, Boolean>() {
    Boolean call(String s) {
        return s.contains("error");
    }
}).count();
```

### Standalone Programs

- Python, Scala, & Java

### Interactive Shells

- Python & Scala

### Performance

- Java & Scala are faster due to static typing
- ...but Python is often fine

# PLATFORM

Built for low-latency

# ABSTRACTION MECHANISM

Resilient **D**istributed **D**atasets

# RDDS

*Immutable Collection*

Accepting **transformations** and **actions**


# TRANSFORMATIONS

- map
- filter
- sample
- union/intersection
- groupByKey
- reduceByKey
- join
- ...

# ACTIONS

- `reduce`
- `collect`
- `count`
- `take(n)`
- `saveAsTextFile`
- ...

# THE FIRST EXAMPLE



$$P(\text{hitting circle}) \approx \text{Surface circle} = \frac{\pi}{4}$$



```
import sc._

val N = 10000000

// Generate a sequence of numbers and distribute
val par = parallelize(1 to N)

// Generate a point in 2D unit square
def randomPoint:(Double,Double) = {
    val x = Math.random()
    val y = Math.random()
    (x,y)
}
// Check if a point lies in the unit circle
def inCircle(point:(Double,Double)):Int = {
    if (point._1*point._1 + point._2*point._2 < 1) 1 else 0
}
```

```
// List of hits yes/no
val inCircleList = par map(i => inCircle(randomPoint))

// Return the first 5 elements from the RDD
inCircleList take 5

// Get info about the RDD
inCircleList.toString

// The number of hits
val total = inCircleList.reduce(_+_ )

// Probability of hitting the circle *4 = Pi
val S = 4. * total / N
```

## From the [Spark examples](#) page

```
val count = parallelize(1 to N).map{i =>
  val x = Math.random()
  val y = Math.random()
  if (x*x + y*y < 1) 1 else 0
}.reduce(_ + _)
println("Pi is roughly " + 4.0 * count / N)
```

# Hadoop M/R in Spark

```
// Read a file from Hadoop FS, (e.g. Ulysses / Project Gutenberg)
// and process it similar to Hadoop M/R
val file = textFile("Joyce-Ulysses.txt")

// Convert to an array of words in the text
val words = file.flatMap(_.split(" "))

// Map to (key,value) pairs
val mapped = words map (word => (word,1))

// Sort and group by key,
// Result is of form (key, List(value1, value2, value3, ...))
val grouped = mapped sortByKey() groupByKey()

// The length of the values array yields the amount
val result = grouped map {case (k,vs) => (k,vs.length)}
// But where is the *reduce*?
```

Be careful with *definitions* of **map** and **reduce**!

```
// Read a file from Hadoop FS, (e.g. Ulysses / Project Gutenberg)
// and process it similar to Hadoop M/R
val file = textFile("Joyce-Ulysses.txt")

// Convert to an array of words in the text
val words = file.flatMap(_.split(" "))

// Map to (key,value) pairs
val mapped = words map (word => (word,1))

// Sort and group by key,
// Result is of form (key, List(value1, value2, value3, ...))
val grouped = mapped sortByKey() groupByKey()

// The length of the values array yields the amount
// val result = grouped map {case (k,vs) => (k,vs.length)}
val result = grouped map {case (k,vs) => (k, vs reduce (_+_))}
```

In Spark, we would use:

```
val file = textFile("Joyce-Ulysses.txt")
val words = file.flatMap(_.split(" "))
val mapped = words map (word => (word,1))
val result = mapped reduceByKey(_+_ )
result collect
```

**... AND REPL**

**... AND WEB INTERFACE**





## Spark Master at spark://ly-1-00:7077

**URL:** spark://ly-1-00:7077

**Workers:** 4

**Cores:** 96 Total, 96 Used

**Memory:** 373.7 GB Total, 128.0 GB Used

**Applications:** 1 Running, 5 Completed

**Drivers:** 0 Running, 0 Completed

### Workers

Id	Address	State	Cores	Memory
<a href="#">worker-20140626200902-ly-1-00.exascience.org-47436</a>	ly-1-00.exascience.org:47436	ALIVE	24 (24 Used)	93.4 GB (32.0 GB Used)
<a href="#">worker-20140626200903-ly-1-01.exascience.org-60945</a>	ly-1-01.exascience.org:60945	ALIVE	24 (24 Used)	93.4 GB (32.0 GB Used)
<a href="#">worker-20140626200903-ly-1-09.exascience.org-48016</a>	ly-1-09.exascience.org:48016	ALIVE	24 (24 Used)	93.4 GB (32.0 GB Used)
<a href="#">worker-20140626200904-ly-2-13.exascience.org-56964</a>	ly-2-13.exascience.org:56964	ALIVE	24 (24 Used)	93.4 GB (32.0 GB Used)

### Running Applications

ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
<a href="#">app-20140627203341-0005</a>	Spark shell	96	32.0 GB	2014/06/27 20:33:41	toniv	RUNNING	49 min

### Completed Applications

ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
<a href="#">app-20140627194920-0004</a>	Spark shell	96	32.0 GB	2014/06/27 19:49:20	toniv	FINISHED	43 min
<a href="#">app-20140626213510-0003</a>	Spark shell	96	32.0 GB	2014/06/26 21:35:10	toniv	FINISHED	14 min
<a href="#">app-20140626211353-0002</a>	Spark shell	96	32.0 GB	2014/06/26 21:13:53	toniv	FINISHED	21 min
<a href="#">app-20140626202327-0001</a>	Spark shell	96	32.0 GB	2014/06/26 20:23:27	toniv	FINISHED	50 min
<a href="#">app-20140626201124-0000</a>	Spark shell	96	32.0 GB	2014/06/26 20:11:24	toniv	FINISHED	11 min

# ... AND DISTRIBUTED MEMORY CACHING

```
val file = textFile("Joyce-Ulysses.txt")
val words = file.flatMap(_.split(" "))
val mapped = words map (word => (word,1))
// Cache the RDD for later use
val cached = mapped cache()
// Use the cached version
val result = cached reduceByKey(_+_ )
// Oops, nothing happens?
result.collect
// Laziness... oh my
result.collect

// Count how many times the word 'the' occurs in the text
cached filter {case(word,v) => word=="the"} reduceByKey(_+_ ) collect
```

**... AND INTERACTIVE USE**

# ... AND THE ECOSYSTEM

- Spark SQL: <http://spark.apache.org/sql/>
- Spark Streaming: <http://spark.apache.org/streaming/>
- BlinkDB: <http://blinkdb.org/>
- MLlib: <http://spark.apache.org/mllib/>
- GraphX: <http://spark.apache.org/graphx/>
- SparkR: <http://amplab-extras.github.io/SparkR-pkg/>

**NOT COVERED**

**BETTER ?**

**FASTER ?**

**EASIER ?**

...

# RDDS UNDER THE HOOD

<http://dl.acm.org/citation.cfm?id=2228301>

# CONFIGURATION & PERFORMANCE

<https://spark.apache.org/docs/latest/configuration.html>

<https://spark.apache.org/docs/latest/tuning.html>




# INSTALLATION & DEPLOYMENT

<https://spark.apache.org/docs/latest/cluster-overview.html>

# EXAMPLE(S)

# GENOMIC DATA



3 billion base pairs ( $3.2 \times 10^9$ )

Packaged in chromosomes

~ 3GB for one human

Analysis requires a lot of processing power and storage

**Transcription Factors:** proteins that bind on region

**Coverage:** basepair occurrence in sequencing

## Coverage data:

Chromosome	Position	Sequencing coverage
19	11004	1
19	11005	2
19	11006	2
19	11007	2
19	11008	3
19	11009	3

# Transcription Factor data:

```
> awk 'BEGIN {srand()} !/^$/ { if (rand() <= .00001) print $0}' bedfile.bed
```

chr1	70529738	70529754	Maf	.	-
chr1	161676477	161676495	Pou2f2	.	-
chr1	176484690	176484699	AP-1	.	-
chr10	6020071	6020084	CTCF	.	-
chr11	1410823	1410838	NF-Y	.	-
chr16	4366053	4366067	YY1	.	+
chr17	77824593	77824602	BAF155	.	+
chr19	10947006	10947013	Rad21	.	-
chr19	49342112	49342121	SIX5	.	+
chr22	39548908	39548922	Irf	.	+
chr7	100048475	100048485	Egr-1	.	-
chr8	119123364	119123374	YY1	.	+
chr8	128562635	128562649	p300	.	-
chr9	14315969	14315982	Egr-1	.	-
chrX	101409366	101409384	CTCF	.	+

```
// Load files from HDFS
val covFile = sc.textFile("NA12878.chrom19.SLX.maq.SRP000032.2009_07.coverage")
val bedFile = sc.textFile("201101_encode_motifs_in_tf_peaks.bed",8)

// Class to hold records from coverage data
class covData(val chr: String, val pos: Int, val cov: Int) {
  def this(line: Array[String]) {
    this(line(0).toString, line(1).toInt, line(2).toInt)
  }
}

// Class to hold records from Transcription Factor data
class tfsData(val chr: String, val pos1: Int, val pos2: Int, val tf: String) {
  def this(line: Array[String]) {
    this(line(0).toString, line(1).toInt, line(2).toInt, line(3).toString)
  }
}
```



```
// Turn input files into an RDD of objects
val cov = covFile.map(_._split("\\s+")).map(new covData(_))
val tfs = bedFile.map(_._split("\\s+")).map(new tfsData(_))

// Count the number of items in both datasets
cov.count
tfs.count

// Cache in memory
val ccov = cov cache
val ctfs = tfs cache

// Count once for the caching to occur
ccov.count
ctfs.count
```

```
// Turn coverage data into K/V pairs
val kvcov = ccov.map(x => (x.pos,(x.cov))).cache
// Turn TF data into K/V pairs
val kvtf = ctfs.filter(x => x.chr == "chr19").map(x => (x.pos1,(x.pos2,x.tf))

// Activate the caching of the coverage data
kvcov.count


// Join both datasets together by key
val cjoined = kvcov.join(kvtf)

// Waaaw, that's fast! In fact, nothing happened yet.
// select 5 entries to see the result but reformat first
val flatjoined = cjoined map { case(x,(y,(z,zz))) => (x,z,zz,y) }
flatjoined take 5

flatjoined.toDebugString
```

# VISUALIZATION OF BIG DATA

# VISUAL ANALYTICS



Visualize information on a chromosome



Create bins and aggregate within bins

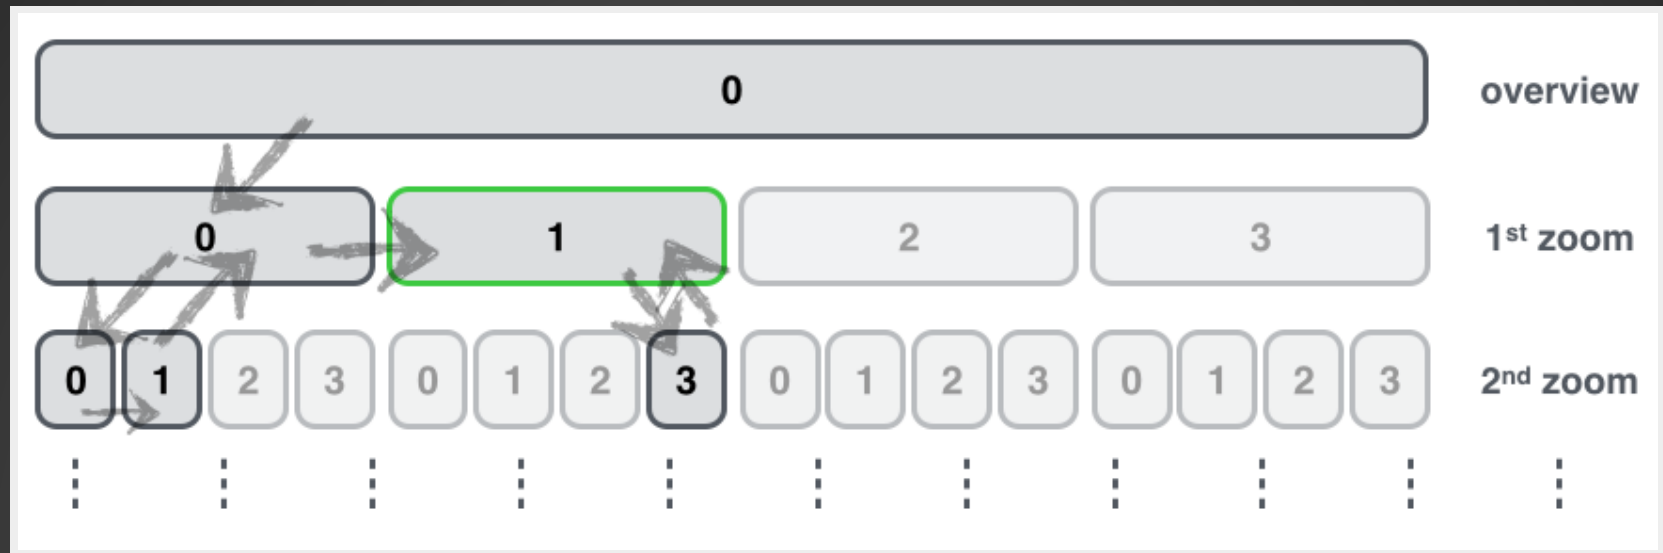














Lazy functional tree zipper

## Part of `treeDraw`:

```
d3.select(window).on("keydown", function() {  
  d3.event.preventDefault();  
  switch (d3.event.keyCode) {  
    case 38:  
      zoomOut(tz,treeDraw);  
      break ;  
    case 40:  
      zoomIn(0,tz,treeDraw);  
      break;  
    case 37:  
      panLeft(tz,treeDraw);  
      break;  
    case 39:  
      panRight(tz,treeDraw);  
      break;  
  };  
});
```

```
object CovQuery extends SparkJob with NamedRddSupport {

  type Range = (Int, Int)
  type TreePath = List[Int]

  val B = 100      // number of bins, also N
  val Z = 10       // zoom level with every step

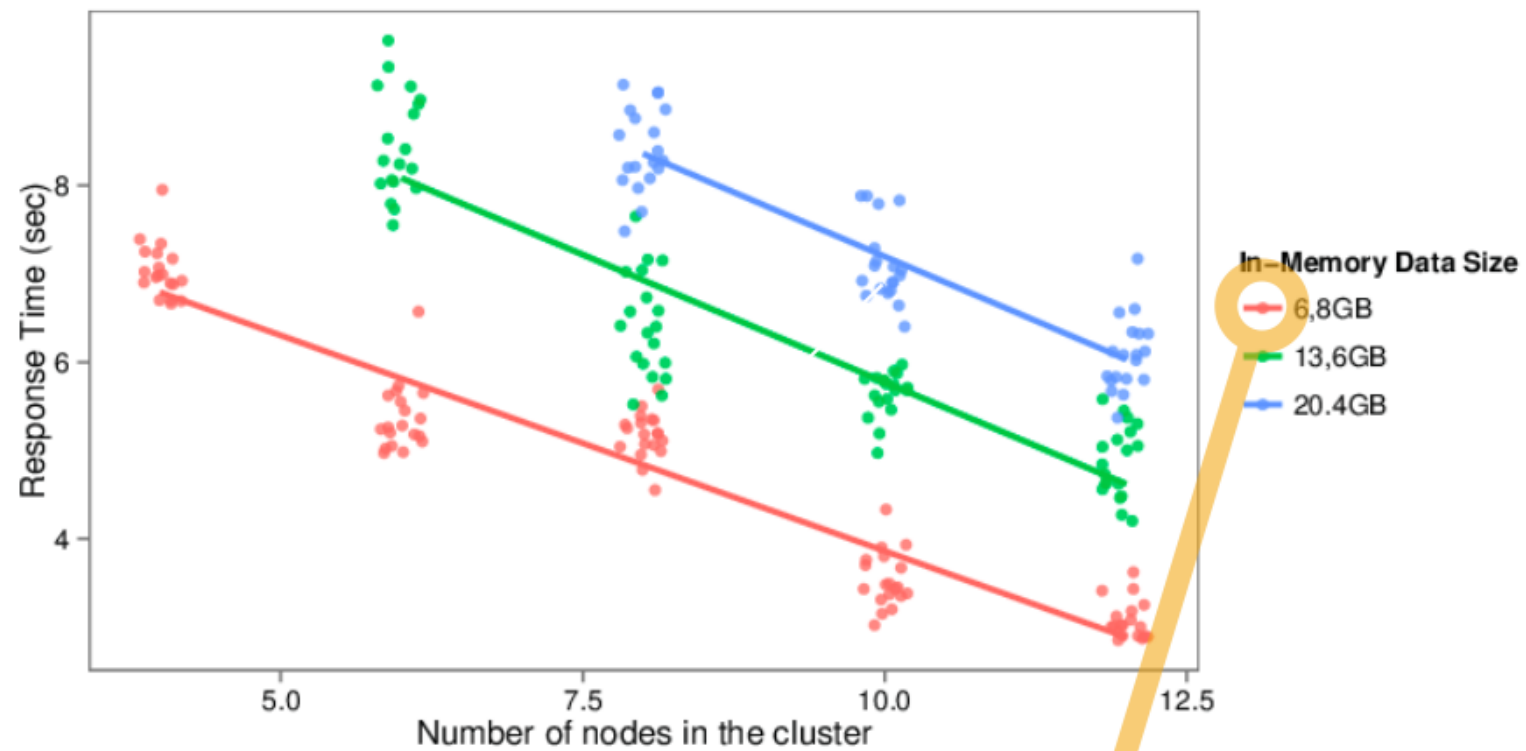
  override def validate(sc: SparkContext, config: Config): SparkJobValidation

  override def runJob(sc: SparkContext, config: Config): Any = {

    // a map with the persistent RDDs
    val listOfPersistentRDDs = sc.getPersistentRDDs

    val chrCache = namedRdds.get[DataPoint](myRDD).get

    // Info about the interval
```



275 million datapoints

# Paper submitted to LDAV 2014

Online Submission ID: 109

## An Architecture For Interactive Big Data Visualization

Category: Research

### ABSTRACT

We present an architecture that allows for interactive visualization of large amounts of data. Visualization of such data often requires pre-processing by aggregating over intervals. Rather than pre-processing the data, our architecture employs a scaling out approach such that adding compute resources has a direct effect on aggregation and visualization performance. We describe a proof-of-concept implementation of this architecture using readily available Open Source technologies. Scaling parameters are derived in order to show that scaling out, in a way transparent to the visualization expert, is a viable option for big data visualization.

**Index Terms:** Computer Graphics [I.3.3]: Picture/Image Generation—Display algorithms Computer Graphics [I.3.3]: Picture/Image Generation—Distributed/network graphics Methodology and Techniques [I.3.6]: Graphics data structures and data types— [Hardware Architecture]: I.3.1—Parallel processing Life and Medical Sciences [J.3]: Biology and genetics—

Another approach to the interactive zooming and panning in large amounts of data is by pre-processing the data and converting it to an in-memory or on-disk representation that can be rendered sufficiently fast [23, 15, 16]. Pre-processing data, e.g. generating data for different zoom levels, requires significantly more storage space than the original data in itself. Additionally, when considering realtime streaming data, online processing is required. In that case, pre-processing and storing intermediate data is not feasible at all.

### 1.3 Aggregation

The data visualization mantra “*overview first, zoom and filter, then details-on-demand*” suggests the requirement for different levels of data aggregation [25]. Since we are bounded by a limited display size, at the overview level, the complete data set is aggregated to fit into one view. The aggregation is a function of the underlying data: sum, maximum/minimum, count, value range, average, etc. Additional zoom levels apply similar aggregation functions to smaller and smaller subsets of the data. This means that by zooming in, the

# THE END

Some links:

- @tverbeiren
- Slides: <https://github.com/tverbeiren/BigDataBe-Spark>
- Spark Home: <https://spark.apache.org/>
- Data Visualization Lab: <http://datavislab.org>
- ExaScience Life Lab: <http://www.exascience.com/>
- Data Intuitive: <http://data-intuitive.com>